

Using `cvc::nav` — and how it maps to GRL-SNAM

Just want the commands to run the demos? See the one-page [cvc-nav-and-grl-snam-quickstart.md](#) — install cvcpkg, run the published Austin demos, or build them yourself. This document is the reference behind it.

A team guide to libcvc's `cvc::nav` module: what it is, how each piece corresponds to the Python GRL-SNAM code you already know, and how to actually use it from either side.

Audience: anyone who works on GRL-SNAM navigation, or who wants to run that navigation inside a C++ host. This covers the base geometry-navigation layer, plus the material-aware extension built on top of it (§9).

Reality check — 2026-09-09. One thing changed since this guide was first written that matters for how you read it: **the obstacle barrier force is now the true gradient of its barrier — fully landed, not just merged.** Earlier the returned derivative was off by $(d - \hat{d}) + 1$, which put an *attraction band* over $(0.39-0.80) \cdot \hat{d}$. Both halves are in: Python (GRL-SNAM#61) and C++ (libcvc#344), **republished as libcvc 3.3.0+cvc.5** so every pycvc that resolves it gets the fix. The coefficient identities below ($\alpha / \beta / \gamma$ = barrier/goal-spring/damping) are unchanged; what changed is that the barrier term is now conservative — and steeper, which forced a `coef_train` learning-rate re-tune (`lr=5e-5`, was `2e-4`). **Every weight trained before this fix is stale.**

This guide remains the base geometry-navigation reference — nothing above touches `cvc::nav`'s own module map in §3.

Contents

- 1. The one-paragraph mental model
- 2. The single most important thing: the fidelity contract
 - A note on the vocabulary: this is not a Hamiltonian system
- 3. The module map (the core of "how it maps")
- 4. The pipeline, stage by stage
 - What "bicycle rollout" actually models
- 5. Getting it running locally
 - Step 0 — install cvcpkg itself
 - Tier 1 — install the runnables (no build at all)
 - Activate the prefix — then everything below is unqualified
 - Tier 2 — build the native kernels (still via cvcpkg)
 - Two canonical builds worth having in your fingers
 - Verify you actually got it
 - Prove the port is faithful (the parity suite)
 - Run the pure-C++ demos
- 6. Using it — Path A: from Python (GRL-SNAM)
 - Feature flags
 - Probing what your pycvc actually carries
 - Running the whole shared-belief navigation runtime in C++ from Python

- 7. Using it — Path B: pure C++ (embed the runtime)
- 8. The learned policy: `.cvcnav` weights
- 9. Beyond geometry: material-aware navigation
- 10. Gotchas & FAQ
- 11. See also

Repos referenced. Two, cloned side by side into a directory of your choosing (`pycvc` is built out of `libcvc` , not a separate repo):

```
git clone https://github.com/CVC-Lab/GRL-SNAM.git
git clone https://github.com/transfix/libcvc.git
```

Every path below is written repo-relative — `libcvc/inc/cvc/nav/drive.h` , `GRL-SNAM/grl_snam/nav.py` — so it resolves wherever you put those two clones. §5 walks through getting the whole thing running from a cold start.

1. The one-paragraph mental model

GRL-SNAM's navigation math is **canonical in Python** (numpy + PyTorch). `cvc::nav` is a **torch-free C++ re-implementation of the navigation hot path** — the distance transform, A^* , the SDF field, the sensor/belief update, the learned coefficient policy, the bicycle drive, and the reactive navigation loop that ties them together. It exists for two reasons:

1. **Speed from Python.** The kernels are a *bit-identical* drop-in that GRL-SNAM calls transparently through `pycvc` . The profiled baseline — 8 agents, real Austin, 384^2 grid, fog on — ran `Squad.step()` at **1375 ms/tick (0.73 Hz)**, where 30 Hz needs 33.3 ms, a $41\times$ gap; the ported kernels close most of it (EDT $43\text{--}65\times$, A $58\text{--}77\times$, *both bit-identical*). See `GRL-SNAM/docs/PERFORMANCE.md` — and note those are that* workload's numbers, not a promise at arbitrary agent counts.
2. **No Python at all.** A renderer or game engine can `#include <cvc/nav/...>` and run the whole navigation loop — sense, plan, drive — with zero libtorch and zero Python.

Python stays the reference twin and the trainer of record; C++ is a fidelity-graded port that never silently redefines the reference. **The status: the port is complete (P0–P8 + CUDA + a torch-free trainer); see `GRL-SNAM/docs/CVCNAV_CPP_PORT_ROADMAP.md` .**

2. The single most important thing: the fidelity contract

Not every part of `cvc::nav` is "the same" as Python in the same sense. The port draws a hard line **at the bilinear SDF sample**. Everything upstream of it is reproduced to the byte; everything from it onward is reproduced to ~ 1 ULP. Know which tier you are standing on before you compare outputs.

Tier	What it covers	What "equal to Python" means	Can it be a transparent default?
BIT-identical	EDT, build_sdf, inflate, line-of-sight, nearest_free, A*, string-pull, sense_batch, to_occupancy / composite / world_to_cell, neighbours	array_equal and .tobytes() match numpy	Yes — this is why the kernels are on by default
FLOAT-equivalent (~1e-4 / ~1 ULP)	sdf_sample, coef_mlp.forward, bicycle_rollout, fused drive_step	allclose(rtol=1e-4... 1e-6) vs torch	No — opt-in forever. torch is the golden generator
BEHAVIORAL	whole-drive sim_world.step vs Swarm.step	same reach-set, no clearance regression, mode-flip rate under budget	No — opt-in
torch-independent	the trainer (coef_train)	finite-difference gradcheck (analytic == numeric)	N/A — it needs no torch reference

Why the line matters. The drive contains a discrete carrot state machine (`stall > 70` , `moved < 0.15` , `dg < best - 1e-3`). A sub-ULP difference in one sample can flip one of those thresholds on a *different* tick and send an agent a different — but equally valid — way. So a single tight position tolerance over a long horizon is the *wrong* test for the drive; that's why the whole-drive tier is "behavioral," and why the float-equivalent drive is never wired transparently into the Python torch path. If you need per-op tightness, compare component-by-component (sampler, MLP, rollout) in lockstep, not end-to-end trajectories.

A note on the vocabulary: this is not a Hamiltonian system

GRL-SNAM's own literature describes the navigator in Hamiltonian-flavored terms, and it's worth being precise about what that does and doesn't buy you in the code this guide covers, so you don't go looking for a conserved quantity that was never implemented.

The published foundational paper does have a genuine port-Hamiltonian formulation — but for exactly one of its three modules (the pose/frame module; the sensor and shape modules are plain symplectic, no port). Neither the GRL-SNAM Python nor `cvc::nav`'s C++ carries that structure forward: there is no `H` defined, evaluated, or differentiated anywhere in either tree; `mass` is a scalar constant used once, as a divisor; there is no momentum state, no port, and nothing named or shaped like a symplectic integrator.

What's actually running, accurately described: a learned-coefficient artificial potential field with viscous damping — the (α, β, γ) triple this guide documents in §3/§8 — advanced by a semi-implicit-Euler-*shaped* update, with the shipped kinematic-bicycle/carrot controller (§4) as the deployment layer on top: pure-pursuit steering plus governors (corner-speed cap, stopping distance, creep floor, reverse), not a potential-field integrator itself. That's a real, working, measured system — it's just not a Hamiltonian one in the structural sense the vocabulary implies, and no energy or passivity argument transfers to it.

Two concrete consequences worth carrying, not just vocabulary:

- **The one place this was a real bug, not just imprecise language:** the obstacle barrier force used to not be the gradient of the barrier it names — see the reality-check box at the top of this guide. That's fixed and republished now, but it's the load-bearing example of why "not actually Hamiltonian" matters operationally and not just semantically: a force field that isn't conservative can't be reasoned about the way the vocabulary suggests.
- **A second, smaller divergence, still open:** `cvc::nav`'s `integrate_surrogate_v2` (`geom_rollout.cpp`) runs **explicit** Euler where its Python namesake runs **semi-implicit** — a genuine cross-language mismatch, though a narrow-blast-radius one (it isn't exported through `pycvc_nav.i` , and its only callers are internal to `geom_rollout.cpp` 's multi-start training loss). Worth knowing before you touch that path; not yet reconciled.

None of this changes how you should use `cvc::nav` day to day — the fidelity contract above is what governs comparing outputs, and it holds regardless of what you call the underlying physics.

3. The module map (the core of "how it maps")

Each `cvc::nav` header, the GRL-SNAM Python it ports, the `pycvc` symbol it's exposed as, the `nav_native` adapter that calls it, and its fidelity tier.

cvc::nav header	GRL-SNAM reference	pycvc.*	nav_nati
grid_nav.h — edt2_squared, build_sdf, inflate, line_of_sight, nearest_free, astar, simplify, *_batch, sense_batch, neighbors_within_radius	grl_snam/planner.py, GRL-SNAM/sdf_nav.py, grl_snam/belief.py (sense)	nav_astar, nav_build_sdf, nav_inflate, nav_line_of_sight, nav_nearest_free, nav_simplify, nav_edt2_squared, nav_*_batch, nav_sense_batch, nav_neighbors	astar, b line_of simplify sense_ba
belief_occupancy.h — to_occupancy, composite_occupancy, world_to_cell	grl_snam/belief.py (BeliefGrid.to_occupancy, composite_occupancy, world_to_cell)	nav_composite_occupancy	to_occup composit
coef_mlp.h — the coef_mlp policy class + .cvcnav I/O	GRL-SNAM/sdf_nav.py CoefMLP + grl_snam/tools/ coef_export.py	nav_coef_mlp_forward	coef_mlp
drive.h — field_stack, sdf_sample, coef_feats, bicycle_rollout, drive_step, carrot_step (+ CUDA)	GRL-SNAM/sdf_nav.py (SDFField.sample, coef_feats, bicycle_rollout); grl_snam/swarm.py (_plan_carrot); grl_snam/ nav.py (SdfNavigator.step)	nav_sdf_sample, nav_coef_feats, nav_bicycle_rollout, nav_drive_step, nav_drive_step_cuda	sdf_samp bicycle_ drive_st
sim_world.h — the reactive navigation runtime	grl_snam/swarm.py Swarm	nav_sim_world_*	NativeSi sim_worl
sim_world_cuda.h — device-resident GPU twin	grl_snam/swarm.py (GPU path)	C++/CUDA only — not bound; Python reaches the GPU via nav_drive_step_cuda	—
sim_thread.h — off-render- thread worker + lock-free snapshot	grl_snam/sim_thread.py	nav_sim_thread_*	NativeSi
coef_train.h — self- supervised trainer (no torch, no labels)	grl_snam/tools/ coef_train.py	nav_train_coef_mlp	train_co
material.h — terrain-risk / hard-hazard navigation	grl_snam/material.py (normative; a port of github.com/SetasAditya/ material-aware-grl-snam)	nav_material_build, nav_material_sample, nav_witness_gate(_batch), nav_bicycle_rollout_material, nav_drive_step_material, nav_integrate_surrogate_material (C++ sim_world::set_material is not bound)	material material witness_ drive_st integrat material
coef_energy_net.h, material_train.h, geom_rollout.h	GRL-SNAM material learned- coefficient net + its training/ surrogate	nav_matnet_forward	matnet_f

Read the right Python reference for the carrot FSM. `grl_snam/nav.py` (`SdfNavigator`, single-agent, list-pop history) is the readable reference, but the C++ `carrot_step/sim_world` is ported from `grl_snam/swarm.py._plan_carrot` (vectorized, ring-buffer history) — that's the twin the parity tests pin. They agree in behavior; they differ in exactly how `moved` is measured.

4. The pipeline, stage by stage

One tick of the reactive drive, and where each stage lives on both sides. This is the data flow the whole module is organized around:

```
flowchart TD
    T["Truth grid"] -->|sense_batch| BF["Log-odds belief (BeliefGrid, per agent)"]
    BF -->|to_occupancy / composite_occupancy| OC["Occupancy raster"]
    OC -->|build_sdf + edt2_squared| SDF["SDF field (phi, normal)"]
    SDF -->|bilinear sample: sdf_sample| SM["Sampled field"]
    SM --> CF["coef_feats (5 features)"]
    SM --> CARROT["Carrot FSM (seek / wall-follow)"]
    GOAL["Goal"] --> CF
    GOAL --> CARROT
    CF --> MLP["CoefMLP.forward (.cvcnav)"]
    MLP --> COEF["(alpha, beta, gamma): barrier / spring / damping"]
    CARROT --> BR["Bicycle rollout: kinematic Ackermann car (nsub substeps)"]
    COEF --> BR
    BR --> POSE["New pose (o, theta, speed)"]
```

Stage	C++ (<code>cvc::nav</code>)	GRL-SNAM
Sense the world into belief	<code>grid_nav.h::sense_batch</code>	<code>belief.py::BeliefGrid.sense</code>
Belief → planning occupancy	<code>belief_occupancy.h::to_occupancy / composite_occupancy</code>	<code>belief.py::BeliefGrid.to_occupancy</code>
Occupancy → signed distance field	<code>grid_nav.h::build_sdf (+ edt2_squared)</code>	<code>sdf_nav.py::build_sdf</code>
Sample the field at the agent	<code>drive.h::sdf_sample (field_stack)</code>	<code>sdf_nav.py::SDFField.sample</code>
Build policy features	<code>drive.h::coef_feats</code>	<code>sdf_nav.py::coef_feats</code>
Policy → coefficients	<code>coef_mlp.h::coef_mlp::forward</code>	<code>sdf_nav.py::CoefMLP.forward</code>
Place the steering carrot (FSM)	<code>drive.h::carrot_step</code>	<code>swarm.py::_plan_carrot / nav.py::_plan_carrot</code>
Integrate one drive tick	<code>drive.h::bicycle_rollout (or fused drive_step)</code>	<code>sdf_nav.py::bicycle_rollout</code>
Whole navigation loop	<code>sim_world.h::sim_world::step</code>	<code>swarm.py::Swarm.step</code>
Route spine over waypoints (A* + lookahead)	<i>(not ported — stays Python)</i>	<code>nav.py::SdfNavigator.drive_to_goal , squad.py::Squad</code>

Note the last row: the multi-waypoint route driver (*A belief spine + lookahead subgoal + closest-approach bookkeeping*) was *deliberately not** ported. It's a thin orchestration layer above `sim_world` that reuses the already-ported `astar / simplify` kernels; if you need it in C++ today, drive it from Python or compose it yourself over `sim_world.retarget`.

What "bicycle rollout" actually models

It is a car, not a bike. The *bicycle model* (a.k.a. the single-track model) is the standard kinematic model of an **Ackermann-steered, car-like vehicle**: the two front wheels are collapsed into one virtual wheel on the centerline, and the two rear wheels likewise, because for the purpose of the path only the centerline matters. The name describes the two-wheel *abstraction*, not the vehicle being simulated. So if you are thinking "Ackermann car," you are thinking of the right thing — with two caveats below.

Both implementations integrate the rear-axle-referenced form:

$$\dot{\theta} = (v / L) \cdot \tan \delta \quad \dot{x} = v \cdot \cos \theta \quad \dot{y} = v \cdot \sin \theta$$

with the vehicle constants living in `veh_params` (C++) / the `bicycle_rollout` keywords (Python) — identical defaults on both sides:

symbol	default	meaning
<code>L</code>	<code>0.035</code>	wheelbase, in the normalized frame (see §10 on coordinates)
<code>delta_max</code>	<code>0.6 rad</code>	steer limit — fixes the minimum turning radius $R_{\min} = L / \tan \delta_{\max}$
<code>vmax</code>	<code>0.9</code>	forward speed cap
<code>a_max</code>	<code>1.5</code>	longitudinal acceleration limit
<code>a_lat_max</code>	<code>1.0</code>	lateral-acceleration cap — this is what makes it slow for corners
<code>k_steer</code>	<code>0.8</code>	how hard the wall barrier is allowed to bias the steer angle

Two things it deliberately does **not** do, worth knowing before you compare it against a vehicle-dynamics model:

- **It does not implement Ackermann steering geometry.** Real Ackermann linkage steers the inner and outer front wheels to different angles so both roll about a common turn centre. Here δ is the single virtual centreline angle (the average of those two); the left/right differential is exactly what the single-track abstraction throws away. The *constraint* an Ackermann car obeys — nonholonomic, bounded turning radius, no strafing — is fully modelled; the *linkage* is not.
- **It is kinematic, not dynamic.** No mass, no tire forces, no slip angle: the rear-axle form above carries no side-slip β term, and cornering speed is bounded by the explicit `a_lat_max` cap rather than emerging from tire saturation.

That is the point, though — a car that cannot strafe is what separates this from the holonomic point-mass rollout (`sdf_rollout`), where a turning radius has no meaning. The learned coefficients (α , β , γ) act through the actuators rather than directly on velocity: the wall barrier *steers* the vehicle away instead of shoving it sideways.

5. Getting it running locally

Everything here runs on your own machine — no lab account, no shared prefix, no special hardware (CUDA is optional throughout). Two tiers, and **you should do tier 1 first**: it is the reference implementation, it needs no C++ toolchain, and it is what the native path is graded against.

Platforms: Linux (x86-64) is fully supported, and **Windows (x86-64) now works** — the Python stack (`pycvc-gl-cp311`, `torch-cp311`, ...) and the native `cvcgl-examples` demos publish for Windows (`.\nav-env\Scripts\activate`). **macOS is coming** — it is blocked on a VTK macOS bundle reaching the catalog. The wasm demos run in any browser. For the terse, copy-pasteable command list across all of these, see the one-page [quickstart](#).

Both tiers go through **cvcpkg**, the group's package manager — that is how everything in this ecosystem is installed and built. Don't hand-roll a venv or call `cmake` directly; the recipes already encode the dependency closure and the build flags, and going around them is how you end up with a mismatched toolchain.

Step 0 — install cvcpkg itself

```
curl -fsSL https://cvcpkg.org/install.sh | sh
```

```
irm https://cvcpkg.org/install.ps1 | iex
```

The bootstrap script detects your platform, downloads a prebuilt binary, verifies its SHA256 and drops it in `$HOME/.local/bin` (override with `CVCPKG_INSTALL_DIR`). It needs no Python and no compiler. If you would rather manage it with pip — or you are on a platform the script does not carry a binary for — `pip install cvcpkg` works too.

```
cvcpkg --version      # 2.0.2 at the time of writing
```

If that fails after a fresh install, `$HOME/.local/bin` is probably not on your `PATH` yet. Full docs: <https://cvcpkg.org/guide>.

Tier 1 — install the runnables (no build at all)

One command gets you a working environment. The `grl-snam-cp31X` column pulls its whole closure — the `pycvc` bindings, numpy, torch — into the prefix:

```
cvcpkg install grl-snam-cp312 --prefix ./nav-env    # or -cp311 / -cp313
```

Activate the prefix — then everything below is unqualified

Every cvcpkg prefix ships an **activation script**, the same ergonomics as a venv. Source it once and your shell is pointed at that prefix:

```
. ./nav-env/bin/activate      # also: activate.csh, activate.fish
```

```
.\nav-env\Scripts\Activate.ps1    # Windows (or Scripts\activate.bat for cmd)
```

It prepends the prefix to everything that matters and saves the old values so it is reversible — run `cvcpkg_deactivate` to put your shell back:

variable	what it gets
<code>PATH</code>	<code><prefix>/bin</code> first — so <code>python</code> , <code>grl-snam</code> and <code>cvc</code> resolve to <i>this</i> prefix
<code>LD_LIBRARY_PATH</code>	<code><prefix>/lib</code> , <code><prefix>/lib64</code> first
<code>CMAKE_PREFIX_PATH</code>	the prefix, so <code>find_package(cvc CONFIG)</code> resolves
<code>PKG_CONFIG_PATH</code>	the prefix's <code>lib/</code> , <code>lib64/</code> , <code>share/</code> <code>pkgconfig</code> dirs
<code>CVCPKG_ACTIVE_PREFIX</code>	the prefix path, for your own scripts to read

Two things worth calling out. It does **not** set `PYTHONPATH`, and it does not need to — putting `<prefix>/bin` on `PATH` means `python` is the prefix's interpreter, which finds `grl-snam` and `pycvc` in its own site-packages. And that `LD_LIBRARY_PATH` ordering is the thing that keeps tier 2 honest: it is the same "fresh lib dir must come first" rule the native build depends on, applied for you rather than exported by hand.

With the prefix active, the commands are just the tools:

```
grl-snam selftest
python -c "import grl_snam, pycvc; print('ok')"
```

`grl-snam` is the unified CLI and the fastest way to confirm a good install — `selftest` needs no data or GPU. Other subcommands worth knowing: `build-sdf`, `train`, `capture`, `fog`, `lab-demo`, `finale`. Add the C++ command-line tool to the same prefix if you want it:

```
cvcpkg install cvc-cli --prefix ./nav-env          # the `cvc` executable
```

That is the full GRL-SNAM reference twin — navigation math, the multi-agent runtime, the trainer — running in numpy/PyTorch. **None of it needs pycvc to carry nav.** With no native kernels present the adapter reports `AVAILABLE == False` and every call transparently takes the Python path, so nothing is gated behind a C++ build.

Heads-up on the packaged bindings. Everything in the catalog is still **3.2.4**, and `cvc::nav` first ships in **3.3.0** — so a `cvcpkg install` today gets you a working GRL-SNAM whose `nav_native.AVAILABLE` is `False`. That is not a broken install; it is the pure-Python path, and it is correct. For the native kernels, build tier 2.

Tier 2 — build the native kernels (still via cvcpkg)

`cvc::nav` is merged into `libcvc master` (3.3.0). The `pycvc-cp31x` recipe takes its source from the checkout it sits in (`source.type: vendored`), so building it out of a master checkout is what produces nav-carrying bindings:

```
cd libcvc
cvcpkg install-deps cvcpkg/recipes/pycvc-cp312 --prefix /path/to/nav-env --config release
cvcpkg build pycvc-cp312 --recipes-dir cvcpkg/recipes --no-deps \
  --prefix /path/to/nav-env --config release
```

Build into the same prefix you activated. A prefix accumulates components, so pointing tier 2 at the tier-1 prefix is what puts the nav-carrying bindings where your activated `python` will find them. Build into a different prefix and the verify below stays `False` with nothing obviously wrong — that mismatch is the single most common way this goes sideways.

How recipes are addressed. Both subcommands take *either* a path to a recipe directory *or* a bare recipe name — a bare name resolves against the default recipe set plus anything you pass with `--recipes-dir`. `cvcpkg` also auto-overlays a `./recipes` directory at the working directory, but **libcvc and GRL-SNAM keep theirs under `cvcpkg/recipes`**, so that auto-overlay never fires for them: pass `--recipes-dir cvcpkg/recipes` (or give the path) or the name will not resolve.

Two more things that bite:

- **`--config` belongs on both commands.** `install-deps` picks the dependency build config and `build` picks yours; mismatch them and you link a release binary against debug deps.
`--no-deps` on `build` is the default, but the project writes it explicitly to make "reuse what install-deps just put here" the visible intent.
- **The recipe still says `upstream_version: 3.2.4`** while master is 3.3.0, so the bundle you build is labelled with the older version even though it carries the nav code. Harmless locally; bump `cvc_revision` before publishing one.

Two canonical builds worth having in your fingers

Both follow the same shape libcvc's own README and CI use: **`install-deps` to populate the prefix, then `build ... --no-deps` to build this checkout into it.** Host tools (cmake, ninja) come from your system `PATH` — add `--include-host-tools` to have `cvcpkg` supply those too.

A. libcvc + the `cvc` CLI, from source, into a prefix. `cvc-cli`'s recipe is `source.type: vendored` (it builds the checkout it sits in) and depends on `libcvc` + `boost`. Order matters: pull the CLI's dependency closure first — that lands a *published* libcvc — then rebuild libcvc from this checkout over it, so the CLI links against your tree rather than the 3.2.4 bundle:

```
cd libcvc
cvcpkg install-deps cvcpkg/recipes/cvc-cli --prefix deps --config release
cvcpkg build libcvc --recipes-dir cvcpkg/recipes --no-deps --prefix deps --config
release
cvcpkg build cvc-cli --recipes-dir cvcpkg/recipes --no-deps --prefix deps --config
release
```

If you only want the SDK and not the CLI, the middle line on its own is the whole build — `install-deps cvcpkg/recipes/libcvc` then `build libcvc`.

B. Pure-Python GRL-SNAM into a prefix. It is published, so the one-liner is just an install — this is what tier 1 above does:

```
cvcpkg install grl-snam-cp312 --prefix deps # or -cp311 / -cp313
```

To go through the recipe instead:

```
cd GRL-SNAM
cvcpkg install-deps cvcpkg/recipes/grl-snam-cp312 --prefix deps --config release
cvcpkg build grl-snam-cp312 --recipes-dir cvcpkg/recipes --no-deps --prefix deps
```

Like the libcvc recipes, `grl-snam-cp31X` is `source.type: vendored` — it builds the checkout it sits in, so `cvcpkg build` picks up your local edits. It is a noarch pure-Python column, so the artifact is valid on every platform.

If your prefix predates GRL-SNAM revision 4, that was not true: revisions 1–3 pinned `python_sdist` to the published `v0.1.0` tarball, so `cvcpkg build` re-built *that release* and silently ignored your tree. Check `source.type` in `cvcpkg/recipes/grl-snam-cp31X/recipe.yaml` if a build mysteriously does not reflect an edit. Either way, for rapid iteration `pip install -e .` under the activated prefix is the shorter loop.

Verify you actually got it

With the prefix activated, a bare `python` is already the right one:

```
python -c "from grl_snam import nav_native; print(nav_native.AVAILABLE,
nav_native.enabled())"
```

`True True` means the bit-identical kernels are live. `False` means you are on the Python path — expected after tier 1, and the sign that a tier-2 build did not land in the prefix you activated. (If you skipped activation, this is the first place it shows up: a system `python` will not see the prefix at all.) If a specific `HAS_*` probe is `False` after a build that otherwise worked, the SWIG wrapper did not regenerate: `touch bindings/pycvc/pycvc_nav.i`, rebuild, and confirm with `hasattr(pycvc, "nav_train_coef_mlp")`.

Prove the port is faithful (the parity suite)

This is the part worth running yourself — it is what makes the fidelity contract in §2 checkable rather than a claim:

```
cd GRL-SNAM
pytest tests/ -k parity -q
```

These compare the C++ kernels against the numpy/torch reference on randomized inputs, tier by tier — `test_nav_cpp_parity.py` (BIT), `test_sdf_sample_parity.py` and `test_coef_mlp_parity.py` (FLOAT), `test_sim_world_parity.py` (BEHAVIORAL). The goldens are generated by forcing the Python backend, so the suite is self-contained.

Every parity test skips itself when pycvc has no nav — so a green run on a tier-1-only machine means "not exercised," not "verified." Check that the tests actually ran (`pytest tests/ -k parity -v`) before concluding anything.

The C++ side has its own tests, which need neither Python nor `pycvc` — ten `nav_*_test` targets that `cvcpkg build` has already produced in its build tree:

```
cd libcvc && ctest --test-dir build -R nav --output-on-failure
```

Run the native demos

The visual `cvc::nav` demos — the reactive swarm on a city, the fog-of-war "ghost" story, the pursuit finale — are the `cvcgl-examples` package, published **prebuilt**. They are interactive, windowed programs; they need no Python and no build:

```
cvcpkg install cvcgl-examples --prefix ./demos
./demos/bin/nav_city_swarm      # also nav_fog_ghost, nav_finale, nav_city_drive,
lsystem_forest
```

Every demo runs with zero args on a synthetic city; point the swarm/drive ones at the real Austin scene with `CVC_NAV_BUNDLE + CVC_NAV_WEIGHTS` (or `--bundle`). To build them yourself: `cvcpkg build cvcgl-examples`, or a direct CMake build of `src/cvcGL` with `-DCVC_BUILD_EXAMPLES=ON` against `cvcpkg install-deps cvcpkg/recipes/cvcgl-examples` (libcvc + VTK + boost + imgui — no Python). The one-page [quickstart](https://transfix.github.io/libcvc/) has the exact flags, and the wasm builds of these demos are live at <https://transfix.github.io/libcvc/>.

Separately, `libcvc/examples/nav_swarm_demo.cpp` is a minimal **headless code template** — no window; it prints pose stats — showing the bare `sim_world` compute loop to copy into your own renderer. It sits behind an off-by-default option:

```
cd libcvc
cvcpkg install-deps cvcpkg/recipes/libcvc --prefix /path/to/nav-env --config release
# with the prefix activated, CMAKE_PREFIX_PATH already points at it
cmake -B build -DCVC_BUILD_NAV_EXAMPLE=ON
cmake --build build -j --target nav_swarm_demo
./build/nav_swarm_demo      # also: nav_material_demo, nav_train_demo
```

`nav_swarm_demo` drives out of the box with no weights file — see §8 on the shipped biased seed.

6. Using it — Path A: from Python (GRL-SNAM)

You rarely call `cvc::nav` directly from Python; you toggle it and let GRL-SNAM dispatch. The adapter is `grl_snam/nav_native.py`.

Feature flags

Flag	Default	Effect
<code>GRL_SNAM_NAV_BACKEND</code>	<code>native</code>	The bit-identical kernels . <code>python</code> forces the pure-numpy reference (the parity tests use this to get the golden).
<code>GRL_SNAM_NAV_DRIVE</code>	<code>torch</code>	The float-equivalent drive inside <code>Swarm</code> . Set <code>native</code> to drive one tick through the C++ <code>sample</code> →MLP→bicycle path.
<code>GRL_SNAM_TRAIN_BACKEND</code>	<code>torch</code>	The trainer . <code>native</code> trains the CoefMLP with the torch-free differentiable rollout. <code>coef_train</code> also takes <code>--backend native</code> , which overrides the env var.

Because the kernels are bit-identical, `native` is the safe default; the drive and trainer are float-equivalent / torch-independent, so they stay explicit opt-ins.

Probing what your `pycvc` actually carries

```
from grl_snam import nav_native

nav_native.AVAILABLE          # pycvc importable AND carries nav_astar
nav_native.enabled()          # AVAILABLE and GRL_SNAM_NAV_BACKEND != "python"

# geometry-navigation layer
nav_native.HAS_SENSE_BATCH    # nav_sense_batch
nav_native.HAS_OCCUPANCY      # nav_composite_occupancy
nav_native.HAS_SDF_SAMPLE     # nav_sdf_sample
nav_native.HAS_COEF_MLP       # nav_coef_mlp_forward
nav_native.HAS_DRIVE          # nav_bicycle_rollout / nav_drive_step
nav_native.HAS_SIM_WORLD      # nav_sim_world_*
nav_native.HAS_SIM_THREAD     # nav_sim_thread_*
nav_native.HAS_CUDA_DRIVE     # nav_drive_step_cuda (pycvc built with CUDA)
nav_native.HAS_TRAIN          # nav_train_coef_mlp

# material-aware layer (§9) – probe these separately; a pycvc can carry the
# geometry kernels and still lack the material ones
nav_native.HAS_MATERIAL              # nav_material_build / _sample / witness gate
nav_native.HAS_MATERIAL_ROLLOUT      # nav_bicycle_rollout_material
nav_native.HAS_MATERIAL_DRIVE        # nav_drive_step_material
nav_native.HAS_MATERIAL_ROLLOUT_INTEGRATOR # nav_integrate_surrogate_material
nav_native.HAS_SIM_WORLD_MATERIAL    # sim_world with material set
nav_native.HAS_MATNET                # the learned material-coefficient net
```

That is the complete set — `dir(nav_native)` filtered on `HAS_` is the authoritative list if it grows.

Every adapter returns the **same Python type** as the function it replaces (a bool ndarray, an `(r, c)` tuple or `None`, a list of `(r, c)`), so a caller dispatches with a one-line early return and nothing downstream can tell which path ran. If `pycvc` is missing or too old, every `HAS_*` is `False`, `enabled()` is `False`, and the Python/torch path runs — importing `grl_snam` is always safe.

Running the whole shared-belief navigation runtime in C++ from Python

```
from grl_snam import nav_native

# Build a native runtime that mirrors a torch Swarm's init (poses, goals,
# colors, field prior, sensor/vehicle params, and the .cvcnav policy):
world = nav_native.sim_world_from_swarm(sw, "coef_mlp.cvcnav", truth=truth_grid)
thread = nav_native.NativeSimThread(world, hz=60.0) # off-thread, GIL-free
thread.start()
frame = thread.read() # (pos_world, heading, speed, mode, reached, tick) or None
thread.retarget(i, gx_n, gy_n)
```

7. Using it — Path B: pure C++ (embed the runtime)

This is the headline capability: thousands of vehicles reacting to a live map with no Python. The whole surface is under `libcvc/inc/cvc/nav/`.

```
#include <cvc/nav/sim_world.h>
using namespace cvc::nav;

sim_world::config cfg;           // rows/cols, world bounds, scale, veh params,
// sensor...
// Zero-setup policy: the shipped biased seed drives out of the box.
// For learned behavior: coef_mlp::load(coef_mlp::default_weights_path())
sim_world world = sim_world::from_occupancy(
    cfg, scene_occupancy, coef_mlp::default_biased(), /*n_agents=*/512);

// Per frame:
world.step();                    // sense (gated) → rebuild → carrot FSM → drive →
// park
world.snapshot(pos_world, heading, speed, mode, reached); // WORLD metres
```

Off the render thread, lock-free:

```
#include <cvc/nav/sim_thread.h>
sim_thread sim(world, /*hz=*/60.0);
sim.start();
auto frame = sim.read();         // whole previous frame or whole next — never torn
sim.retarget(i, gx_n, gy_n);     // live edits flow in through a command queue
```

A compilable template lives at `libcvc/examples/nav_swarm_demo.cpp`, built with `-DCVC_BUILD_NAV_EXAMPLE=ON` (alongside `nav_material_demo.cpp` and `nav_train_demo.cpp`).

Belief modes (the C++ counterpart of `Swarm.belief_mode`), selected by a per-agent `map_id`:

- `shared` ($M = 1$) — one belief plane; the thousands-of-agents deployment path.
- `clustered` (K groups) — one plane per group; isolation is structural.
- `private_belief` ($M = N$) — one plane per agent; the fog-of-war fidelity twin. (Spelled `private_belief`, not `private` — the latter is a C++ keyword.)

Zero-setup ergonomics worth knowing: `coef_mlp::default_biased()` (drive with no weights file — the constant $(\alpha, \beta, \gamma) = (1, 3, 4)$ basin), `sim_world::from_occupancy` (scatter N agents on free cells of any rasterized map), and `coef_mlp::default_weights_path()` (resolve the installed policy).

8. The learned policy: `.cvcnav` weights

The CoefMLP is a tiny $5 \rightarrow 64 \rightarrow 64 \rightarrow 3$ SiLU net whose output is `softplus(net(feat) + log(expm1(bias)))` — the $(\alpha, \beta, \gamma) = \text{barrier} / \text{goal-spring} / \text{damping}$ the rollout uses. One versioned binary blob feeds torch, the C++ CPU forward, and the CUDA forward alike.

- **Export** a trained torch checkpoint with `grl_snam/tools/coef_export.py` (`python -m grl_snam.tools.coef_export coef_sdf.pt coef_mlp.cvcnav`).
- **Default location** the C++ resolves, in order: `CVC_NAV_WEIGHTS` env var → the baked `$PREFIX/share/cvc/nav/coef_mlp.cvcnav` → a path relative to the loaded `libcvc` (for relocated / `cvcpkg` installs). That last fallback is `dladdr`-based and compiled out on Windows (`#ifndef _WIN32`), so a relocated Windows install resolves only the first two — set `CVC_NAV_WEIGHTS` there if the baked path moved.
- **arch_hash** in the header pins the architecture: a hidden-size change bumps the file and the loader refuses a mismatched net rather than reading garbage.

Heads-up — the biased seed is currently the best base policy we have. The *installed default* weights are the **untrained biased seed** (~57–61% reach); zero-config driving works, but it is the hand-tuned (1,3,4) basin, not a trained checkpoint. Do **not** assume training beats it: in current measurements, training the base CoefMLP with the self-supervised rollout *regresses* — reach holds near the seed to ~25 steps, then falls to ~0.15–0.23 by 100 steps on **both** the torch `reach_rate` and the native `sim_world` eval paths, so the older “~65% target” does **not** reproduce for the base net. Root cause: the trainer optimizes random free-space surrogate rollouts, not the deployed carrot/route-guided regime. The lever for better demo nav is the **global field/route**, not the CoefMLP weights. Concretely: a **clearance-weighted A* surcharge** — `sdf_nav.clearance_cost` (a hinge `gamma*max(0, d_safe - clearance)` off the SDF clearance field) wired onto each agent's `route_cost_fn` via `squad.attach_clearance_routing` — makes the route trade a little length for obstacle standoff, which the local drive follows more reliably. Measured on the route-guided city squad (which `reach_rate` can't see — it drives a route-less Swarm, n=20 × 5 seeds): reach ~**0.80** → ~**0.90** at (`d_safe=6`, `gamma=1.5`). It is **budget-sensitive** — the standoff route is longer, so give the run tick headroom or the longer routes get cut off and reach *regresses*. Keep the seed installed unless a trained checkpoint actually beats it on your own reach eval (mean up **and** no per-seed regression). See `libcvc/docs/NAV_TRAINING.md`.

Pretrained weights package: `grl-snam-weights`

So you don't have to run a training pass to get a trained (α , β , γ) net, a trained base CoefMLP ships as a **public** cvcpkg data package on the `cvc` org (no token — this is the research tier):

```
cvcpkg install grl-snam-weights --prefix ./nav-env
# installs, under the prefix:
#   share/grl-snam-weights/coef_sdf.cvcnav  # native CVNV blob – cvc::nav forward
#   share/grl-snam-weights/coef_sdf.pt      # torch checkpoint – grl_snam / sdf_nav
#   share/grl-snam-weights/PROVENANCE.md
```

- **C++ (`cvc::nav`)** — `default_weights_path()` (§8) now also probes `share/grl-snam-weights/coef_sdf.cvcnav` as a **last resort**, so an installed package is discovered *if the canonical seed is absent*. But the shipped seed (`share/cvc/nav/coef_mlp.cvcnav`) is always present and **wins by default** — and currently out-performs this trained base net — so to actively use the packaged weights set `CVC_NAV_WEIGHTS=.../share/grl-snam-weights/coef_sdf.cvcnav` (checked first), or copy it over the baked `coef_mlp.cvcnav`.
- **Python (GRL-SNAM)** — `sdf_nav.CoefMLP().load_state_dict(torch.load( ".../coef_sdf.pt")["model_state_dict"] )`.

What this is (and is not). `v1.0.0` is a *proof-of-pipeline* checkpoint — a short self-supervised run on the `austin_south` SDF that demonstrates the train → export → package → pull → load path end to end. It is **not** claimed to beat the installed biased seed: per the heads-up above, training the base CoefMLP on the free-space rollout surrogate currently *regresses* versus the hand-tuned (1,3,4) basin. Use this package to exercise the pipeline and as a starting checkpoint; keep the seed unless a longer run beats it on your own reach eval. The **RF-comm-aware** DBG policy is a *separate, private* package (`cvc-dbg-weights` on the `utdbg` org) — the research/development seam is kept at the org boundary, never mixed into this public package.

9. Beyond geometry: material-aware navigation

`cvc::nav` can navigate over *terrain semantics*, not just walls: a per-cell **risk** field (mud, rubble — soft costs that bias but never block) and a **hard hazard** mask (water, cliffs — lethal but not physical geometry), adding two force terms and a feasibility "witness gate" on top of the SDF drive. It's off by default (byte-unchanged runs) and turned on per-world with `sim_world::set_material(...)`. Header `libcvc/inc/cvc/nav/material.h`; full write-up in `libcvc/docs/NAV_MATERIAL.md`. Python normative reference: `grl_snam/material.py`.

Where the geometry drive reads its (α, β, γ) from the small `CoefMLP` (§8), the material drive's soft/hard weights come from a **learned material-coefficient net**, `CoefEnergyNetMaterial` — a transformer over obstacle/goal tokens plus a CNN risk-patch encoder (`coef_energy_net.h`, weights in the `.cvcnm` format, GRL-SNAM reference `grl_snam`'s `matnet`). It has the full port treatment: a torch-free CPU `forward_batch / backward_one` (finite-difference gradchecked), and a **CUDA twin** — `forward_batch_cuda / backward_batch_cuda` (float-equivalent to the CPU, one block per agent) plus `material_adam_cuda`, a device-resident optimizer — so the whole material training step can run on the GPU. Probe `nav_native.HAS_MATNET`; the device path is `CVC_ENABLE_CUDA`-gated and auto-skips without a device.

10. Gotchas & FAQ

- **The published `pycvc` does not carry nav yet.** Everything on the package index is still `3.2.4`; `cvc::nav` lands in `3.3.0`. So a packaged install reports `nav_native.AVAILABLE == False` and runs the Python path — correct behavior, not a broken install. Build tier 2 in §5 for the native kernels.
 - **Most "it didn't work" reports are an un-activated shell.** `cvcpkg install` puts everything in the prefix and changes nothing about your current shell, so until you source `<prefix>/bin/activate` (or `Scripts\Activate.ps1`) you are still running the system `python` and the system libraries. Check `$CVCPKG_ACTIVE_PREFIX` before debugging anything else.
 - **A green parity run may mean "skipped."** Every parity test skips itself when `pycvc` has no nav, so the suite passes on a tier-1-only machine without exercising a single C++ kernel. Run it with `-v` and confirm the tests actually ran before treating a pass as verification.
 - **`reach_to1` is in normalized units (0.8), not metres.** Both `Swarm` and `SdfNavigator` default `0.8`. Do not "fix" it to `0.15`; that changes the reach/park set on every run.
 - **`nsub` (bicycle substeps) defaults to 1 at deployment.** Raise it only where thin-wall tunnelling matters; at the sense-bound steady state its effect is second-order.
 - **Don't expect the drive to match torch trajectory-for-trajectory.** It's float-equivalent per op; over a long horizon the carrot FSM's threshold decisions make trajectories diverge while both stay valid. Compare component-wise or on the behavioral gate (see §2).
 - **Coordinate regime.** The drive works in a **normalized, centered** frame: `world = normalized / scale + center`. `sim_world` snapshots convert back to world metres for you; the raw sampler/drive functions take normalized positions.
 - **The C++ path never becomes the Python default.** `GRL_SNAM_NAV_DRIVE / GRL_SNAM_TRAIN_BACKEND` are opt-in by design — torch stays the reference and golden generator. Only the bit-identical kernels are on by default.
-

11. See also

- `GRL-SNAM/docs/NATIVE_CVC_NAV.md` — the GRL-SNAM-side reference (flags, `HAS_*`, getting a `pycvc` that carries the layer).
- `GRL-SNAM/docs/CVCNAV_CPP_PORT_ROADMAP.md` — the C++ port design record (phases P0–P8, the fidelity boundary, `.cvcnav` format, `sim_world / sim_thread`).
- `GRL-SNAM/docs/PERFORMANCE.md` — measured kernel speedups and the scaling wall.
- `libcvc/docs/NAV_TRAINING.md` — the torch-free trainer's API.
- `libcvc/docs/NAV_MATERIAL.md` — material-aware navigation.
- Headers: `libcvc/inc/cvc/nav/` · Demos: `libcvc/examples/nav_*.demo.cpp` (- `DCVC_BUILD_NAV_EXAMPLE=ON`).